

Housing dataset

August 31, 2022

```
[1]: # 1. Load housing dataset "housing.csv" (find the file on this server) into a
      ↪variable df
import pandas as pd
df = pd.read_csv(r'/home/ishita/Documents/datasets/housing/housing.csv')
```

```
[3]: # 2. Display the brief information about this dataset.
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
[4]: # 3. Display number of rows and features available in this dataset.
print("Number of rows: ",df.shape[0])
print("Number of coloumn: ",df.shape[1])
```

```
Number of rows: 20640
Number of coloumn: 10
```

```
[5]: # 4. Find the target variable.
target = "median_house_value"
print(df[target])
```

```
0          452600.0
```

```

1      358500.0
2      352100.0
3      341300.0
4      342200.0

```

```

...
20635    78100.0
20636    77100.0
20637    92300.0
20638    84700.0
20639    89400.0

```

Name: median_house_value, Length: 20640, dtype: float64

```
[6]: # 5. Show first few rows of the dataset.
df.head()
```

```
[6]:  longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0    -122.23    37.88             41.0           880.0           129.0
1    -122.22    37.86             21.0          7099.0          1106.0
2    -122.24    37.85             52.0          1467.0           190.0
3    -122.25    37.85             52.0          1274.0           235.0
4    -122.25    37.85             52.0          1627.0           280.0

    population  households  median_income  median_house_value  ocean_proximity
0         322.0        126.0         8.3252         452600.0        NEAR BAY
1        2401.0       1138.0         8.3014         358500.0        NEAR BAY
2         496.0        177.0         7.2574         352100.0        NEAR BAY
3         558.0        219.0         5.6431         341300.0        NEAR BAY
4         565.0        259.0         3.8462         342200.0        NEAR BAY
```

```
[7]: # 6. Display the summary statistics about all the features of the dataset.
df.describe()
```

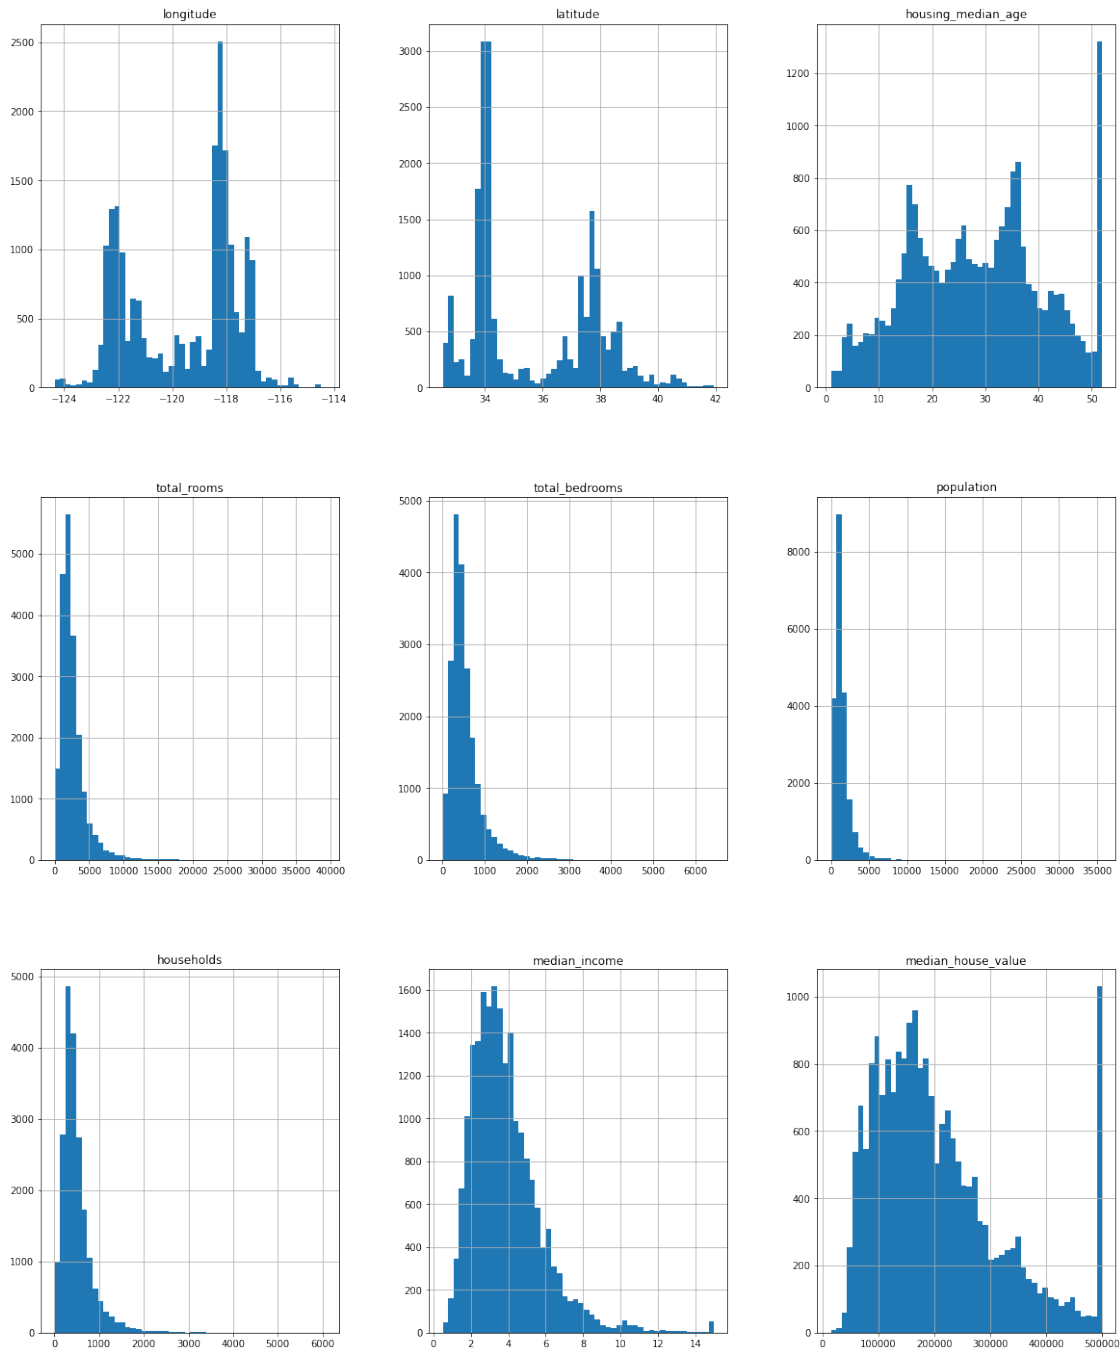
```
[7]:  longitude  latitude  housing_median_age  total_rooms  \
count  20640.000000  20640.000000         20640.000000  20640.000000
mean    -119.569704    35.631861          28.639486    2635.763081
std         2.003532     2.135952         12.585558    2181.615252
min     -124.350000    32.540000          1.000000     2.000000
25%     -121.800000    33.930000         18.000000    1447.750000
50%     -118.490000    34.260000         29.000000    2127.000000
75%     -118.010000    37.710000         37.000000    3148.000000
max     -114.310000    41.950000         52.000000   39320.000000

    total_bedrooms  population  households  median_income  \
count  20433.000000  20640.000000  20640.000000  20640.000000
mean     537.870553   1425.476744    499.539680     3.870671
std     421.385070   1132.462122    382.329753     1.899822
min       1.000000     3.000000     1.000000     0.499900
```

25%	296.000000	787.000000	280.000000	2.563400
50%	435.000000	1166.000000	409.000000	3.534800
75%	647.000000	1725.000000	605.000000	4.743250
max	6445.000000	35682.000000	6082.000000	15.000100

	median_house_value
count	20640.000000
mean	206855.816909
std	115395.615874
min	14999.000000
25%	119600.000000
50%	179700.000000
75%	264725.000000
max	500001.000000

```
[8]: # 7. Show the histogram plot of each attribute.
%matplotlib inline
import matplotlib.pyplot as plt
df.hist(bins=50, figsize=(20,25))
plt.show()
```



```
[13]: # 8. Show if there are any missing/Null values in the dataset.
df.isnull().sum()
```

```
[13]: longitude      0
      latitude      0
      housing_median_age  0
```

```

total_rooms          0
total_bedrooms       207
population            0
households            0
median_income         0
median_house_value    0
ocean_proximity      0
dtype: int64

```

```

[12]: # 9. Show different types of values in categorical attributes along with their
      ↪ frequencies.
      df["ocean_proximity"].value_counts()

```

```

[12]: <1H OCEAN      9136
      INLAND       6551
      NEAR OCEAN   2658
      NEAR BAY     2290
      ISLAND        5
      Name: ocean_proximity, dtype: int64

```

```

[13]: # 10. Fill the missing values with most frequently used value for categorical
      ↪ attribute and for
      # numerical attribute fill median value.
      df_num = pd.read_csv(r'/home/ishita/Documents/datasets/housing/housing.csv')
      median = df_num["total_bedrooms"].median()
      df_num["total_bedrooms"].fillna(median, inplace=True)
      mode = df_num["ocean_proximity"].mode()
      df_num["ocean_proximity"].fillna(mode, inplace=True)
      df_num.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20640 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

```
[14]: # 11. Display sum of missing values after filling the values.
      ((df.isnull().sum().sum()) * median)
```

```
[14]: 90045.0
```


```
[15]: # 12. Transform "median_income" attribute into a new attribute "income_cat"
      ↪which has 5 levels (1,2,3,4,5) ranging
      #from 0-1.5, 1.5-3.0, 3.0-4.5, 4.5-6.0, 6.0-np.inf respectively.
      # Use pd.cut(df["median_income"], bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
      ↪labels=[1, 2, 3, 4, 5])
      import numpy as np
      df["income_cat"] = pd.cut(df["median_income"],
      bins=[0., 1.5, 3.0, 4.5, 6., np.inf],
      labels=[1, 2, 3, 4, 5])
```

```
[16]: # 13. Find the distribution based on "income_cat" in the entire dataset.
      df["income_cat"].value_counts() / len(df)
```

```
[16]: 3    0.350581
      2    0.318847
      4    0.176308
      5    0.114438
      1    0.039826
      Name: income_cat, dtype: float64
```

```
[17]: # 14. Plot histogram of "income_cat" attributes. (use df['attribute name'].
      ↪hist() )
      df["income_cat"].hist()
```

```
[17]: <AxesSubplot:>
```



```
[18]: # 15. Split the dataset 80% of rows for training, and 20% of rows for testing.
      ↳ purpose. Just for the sake
      ↳ of learning take first 80% rows as training and, rest 20% rows as testing
      ↳ respectively. Store these
      ↳ train and test datasets in temp_train and temp_test variables.
      def split_train_test(data, test_ratio):
          shuffled_indices = np.random.permutation(len(data))
          test_set_size = int(len(data) * test_ratio)
          test_indices = shuffled_indices[:test_set_size]
          train_indices = shuffled_indices[test_set_size:]
          return data.iloc[train_indices], data.iloc[test_indices]
      temp_train, temp_test = split_train_test(df, 0.2)
      print(len(temp_train))
      print(len(temp_test))
```

```
16512
4128
```

```
[19]: # 16. Check the distribution based on "income_cat" in train and test set that
      ↳ you obtained in above step.
      df["income_cat"].value_counts() / len(df)
```

```
[19]: 3    0.350581
      2    0.318847
      4    0.176308
      5    0.114438
```

```
1    0.039826
Name: income_cat, dtype: float64
```

```
[20]: temp_train["income_cat"].value_counts() / len(temp_train)
```

```
[20]: 3    0.349746
      2    0.317769
      4    0.177386
      5    0.115250
      1    0.039850
      Name: income_cat, dtype: float64
```

```
[21]: temp_test["income_cat"].value_counts() / len(temp_test)
```

```
[21]: 3    0.353924
      2    0.323159
      4    0.171996
      5    0.111192
      1    0.039729
      Name: income_cat, dtype: float64
```

```
[22]: # 17. Reshuffle the dataset to have stratified distribution of 'income_cat' and
      ↪ then split it into train
      #and test. Use following function
      from sklearn.model_selection import StratifiedShuffleSplit
      split = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=42)
      for train_index, test_index in split.split(df, df["income_cat"]):
          train = df.loc[train_index]
          test = df.loc[test_index]
      print(len(train))
      print(len(test))
```

```
16512
4128
```

```
[23]: # 18. Check again the distribution based on "income_cat" in train and test set
      df["income_cat"].value_counts() / len(df)
```

```
[23]: 3    0.350581
      2    0.318847
      4    0.176308
      5    0.114438
      1    0.039826
      Name: income_cat, dtype: float64
```

```
[24]: train["income_cat"].value_counts() / len(train)
```



```
[24]: 3    0.350594
      2    0.318859
      4    0.176296
      5    0.114462
      1    0.039789
      Name: income_cat, dtype: float64
```

```
[25]: test["income_cat"].value_counts() / len(test)
```

```
[25]: 3    0.350533
      2    0.318798
      4    0.176357
      5    0.114341
      1    0.039971
      Name: income_cat, dtype: float64
```

```
[26]: # 19. Find correlation of target attribute with rest of the attributes. Use
      ↪ correlation=df.corr()
      # correlation["attribute name"].sort_values()
      correlation=df.corr()
      correlation["median_house_value"].sort_values()
```

```
[26]: latitude          -0.144160
      longitude         -0.045967
      population        -0.024650
      total_bedrooms     0.049686
      households         0.065843
      housing_median_age 0.105623
      total_rooms        0.134153
      median_income      0.688075
      median_house_value 1.000000
      Name: median_house_value, dtype: float64
```

```
[27]: # 20. Convert categorical attribute to numeric using ordinal encoder. Use
      from sklearn.preprocessing import OrdinalEncoder
      oe=OrdinalEncoder ()
      df_cat_oe =oe.fit_transform(df[["ocean_proximity"]])
      print(df_cat_oe)
```

```
[[3.]
 [3.]
 [3.]
 ...
 [1.]
 [1.]
 [1.]]
```

```
[28]: # 21. Add the new attribute that you have transformed into numeric into dataset
      ↪ df.
      d = dict(enumerate(df_cat_oe.flatten(), 1))
      df["df_cat_oe"]=d.values()
      df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
10  income_cat             20640 non-null  category
11  df_cat_oe              20640 non-null  float64
dtypes: category(1), float64(10), object(1)
memory usage: 1.8+ MB
```

```
[29]: # 22. Drop the attribute which has categorical values from the dataset.
      df=df.drop("ocean_proximity" , axis=1)
```

```
[30]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population             20640 non-null  float64
6   households              20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   income_cat             20640 non-null  category
10  df_cat_oe              20640 non-null  float64
```

```
dtypes: category(1), float64(10)
memory usage: 1.6 MB
```

```
[54]: # 23. Split the dataset. use sklearn.model_selection import train_test_split
# train_set, test_set = train_test_split(housing, test_size=0.2,
↳ random_state=42)
from sklearn.model_selection import train_test_split
train_set, test_set = train_test_split(df, test_size=0.2, random_state=42)
print(len(train_set))
print(len(test_set))
```

```
16512
4128
```

```
[56]: # 24. Separate the target attribute and rest of the attributes from train_set,
↳ and test_set and store
# them as train_target, and test_target in two separate variables.
train_target = train_set[target]
train_rest=train_set.drop("median_house_value",axis=1)
test_target = test_set[target]
test_rest = test_set.drop("median_house_value",axis=1)
train_target.info()
print("\n")
test_target.info()
```

```
<class 'pandas.core.series.Series'>
Int64Index: 16512 entries, 14196 to 15795
Series name: median_house_value
Non-Null Count  Dtype
-----
16512 non-null  float64
dtypes: float64(1)
memory usage: 258.0 KB
```

```
<class 'pandas.core.series.Series'>
Int64Index: 4128 entries, 20046 to 3665
Series name: median_house_value
Non-Null Count  Dtype
-----
4128 non-null  float64
dtypes: float64(1)
memory usage: 64.5 KB
```

```
[70]: # 25. Take a linear regression mode and train it.
from sklearn.linear_model import LinearRegression
reg = LinearRegression()
```

```
[71]: # 26. reg.fit(training_dataset_name, training_dataset_target)
      reg.fit(train_set, train_target)
      reg.intercept_
      reg.coef_
```

```
[71]: array([-5.53030717e-11, -5.81505955e-11,  3.94919286e-14,  2.15175100e-14,
          -9.03183778e-14, -7.23726634e-15,  2.01839413e-14, -1.20954654e-11,
           1.00000000e+00,  5.77589516e-13, -9.05660909e-14])
```

```
[73]: # 27. Predict few values from the dataset. Use predict method and pass some
      ↪ rows from dataset.
      train_set_rand = train_set.sample(frac=0.50)
      reg.predict(train_set_rand)
```

```
[73]: array([184900., 158200., 328500., ..., 262700., 113600., 180100.])
```